

MIT Sloan
Management Review



BIG IDEAS
RESEARCH REPORT

In collaboration with

BCG

November 2025

The Emerging Agentic Enterprise: How Leaders Must Navigate a New Age of AI

by Sam Ransbotham, David Kiron,
Shervin Khodabandeh, Sesh Iyer,
and Amartya Das

AUTHORS

Sam Ransbotham is a professor of analytics at the Carroll School of Management at Boston College, as well as guest editor for *MIT Sloan Management Review's* Artificial Intelligence and Business Strategy Big Ideas initiative.

David Kiron is the editorial director, research, of *MIT Sloan Management Review* and program lead for its Big Ideas research initiatives.

Shervin Khodabandeh is a managing director and senior partner at Boston Consulting Group (BCG) and the coleader of its AI business in North America. He is a leader in BCG X and has over 20 years of experience driving business impact from AI and digital. He can be contacted at shervin@bcg.com.

Sesh Iyer is a managing director and senior partner at BCG and the North America chair for BCG X, where he helps clients drive large-scale AI transformations. He is the global leader for the AI & Tech Lab at the BCG Henderson Institute and can be reached at sesh@bcg.com.

Amartya Das is a principal at BCG and currently serves as an ambassador at the BCG Henderson Institute, where he leads research on the impact of technology and AI on society, focusing on how emerging technologies reshape both companies and public institutions. He can be reached at das.amartya@bcg.com.

CONTRIBUTORS

Todd Fitz, Kevin Foley, Sarah Johnson, Matt Langione, Michele Lee DeFilippo, Amanda Luther, Jennifer Martin, Samantha Oldroyd, Meenal Pore, Allison Ryder, Taylor Zhang, Leila Zhu, Leonid Zhukov, David Zuluaga Martínez

The research and analysis for this report was conducted under the direction of the authors as part of an *MIT Sloan Management Review* research initiative in collaboration with and sponsored by Boston Consulting Group.

To cite this report, please use:

S. Ransbotham, D. Kiron, S. Khodabandeh, S. Iyer, and A. Das, “The Emerging Agentic Enterprise: How Leaders Must Navigate a New Age of AI,” *MIT Sloan Management Review* and Boston Consulting Group, November 2025.


SUPPORTING SPONSORS



<https://doi.org/10.63383/jAXE2583>

CONTENTS

- 1** Introduction
 - 4** Strategic Tensions When Adopting Agentic AI
 - 11** A Strategic Overhaul of Workflows, Governance, Roles, and Investment
 - 17** Conclusion
 - 18** Appendix
-



Introduction

Executives have long relied on simple categories to frame how technology fits into organizations: Tools automate tasks, people make decisions, and strategy determines how the two work together. That framing is no longer sufficient. A new class of systems — agentic AI — complicates these boundaries. These systems can plan, act, and learn on their own. They are not just tools to be operated or assistants waiting for instructions. Increasingly, they behave like autonomous teammates, capable of executing multistep processes and adapting as they go. Notably, 76% of respondents to our global executive survey say they view agentic AI as more like a coworker than a tool.


For strategists, agentic AI's dual nature as both a tool and coworker creates new dilemmas. A single agent might take over a routine step, support a human expert with analysis, and collaborate across workflows in ways that shift decision-making authority. This tool-coworker duality breaks down traditional management logic, which assumes that technology either substitutes or complements, automates or augments, is labor or capital, or is a tool or a worker, but not all at once. Organizations now face an unprecedented challenge: managing a single system that demands both human resource approaches and asset management techniques.

The separation of technology and strategy inside most organizations exacerbates this challenge. Technology executives focus on technology issues, making pilot, vendor, or infrastructure decisions. Strategic executives focus on markets, competition, and people. But agentic AI makes that separation untenable. It simultaneously influences the design of processes, the structure of roles, the allocation of decision rights, and the culture of accountability.

ABOUT THE RESEARCH

This report presents findings from the ninth annual global research study on artificial intelligence and business strategy by *MIT Sloan Management Review* and Boston Consulting Group. In spring 2025, we fielded a global survey and subsequently analyzed records from 2,102 respondents representing more than 21 industries and 116 countries. We also interviewed 11 executives leading AI initiatives in a broad range of companies and industries, including financial services, technology, retail, energy, and health care.

Our research examines the speed of agentic AI adoption inside organizations, which has outpaced the adoption of traditional and generative AI. Exploring how agentic AI relates to capital and labor, this report outlines a series of resultant tensions and offers suggestions for organizations seeking to resolve them.



A Tidal Wave of Adoption, a Trickle of Strategy

Despite the technology's wide-ranging implications, organizations are rapidly adopting agentic AI, well before they have a strategy in place. Their adoption of traditional AI has climbed to 72% over the past eight years, according to our survey. (SEE FIGURE 1.) Generative AI achieved 70% adoption in just three years. In just two years, agentic AI has already reached 35% adoption, with another 44% of organizations planning to deploy it soon. Vendors are accelerating this trend by embedding agentic capabilities as features in their offerings, causing organizations to implement agentic AI before they have developed a strategic management framework.

The result is a growing strategic risk: Agentic AI is spreading across enterprises faster than leaders can redesign processes, assign decision rights, or rethink workforce

models. Our research, based on 2,000-plus respondents to a global survey and interviews with leading executives, finds that organizations have multiple options for obtaining value from agentic AI. It offers possibilities not only to improve cost efficiency but also to expand revenue, accelerate innovation, compress learning curves, and restructure organizations. Without a strategic approach that aligns these objectives, organizations risk limiting returns on their investments.

The Executive Dilemma

Agentic AI's dual nature as both a tool and coworker creates competing organizational pressures that management frameworks cannot resolve. IT leaders want predictable, scalable systems with clear technical specifications. CFOs need investment models with measurable returns and depreciation schedules. HR executives require




FIGURE 1
Traditional AI Adoption
Continues to Grow

Since 2023, the percentage of organizations piloting or deploying AI solutions has risen 22 percentage points.

performance management frameworks and supervision protocols. Business leaders demand both efficiency and adaptability from the same system.

These competing demands aren't implementation challenges. They're strategic imperatives that expose new sources of organizational differentiation. Among organizations with extensive agentic AI use, we find that 73% believe using AI fundamentally increases their ability to stand out, while 76% of their employees believe it changes how individuals differentiate themselves from coworkers.

WHY AGENTIC AI SPREADS SO FAST

The rapid adoption of agentic AI aligns perfectly with the diffusion of innovations theory, which explains why some technologies spread faster than others.¹ According to the framework, innovations diffuse more quickly when they offer a clear relative advantage, are compatible with existing systems, are simple to try, and have observable results.

Agentic AI excels on all four dimensions. It builds on the familiarity of generative AI that organizations have already adopted, integrates seamlessly into widely used platforms, requires minimal new infrastructure for experimentation with it, and delivers immediate, visible productivity benefits. As Chevron's chief data and analytics officer, Margery Connor, explains, "With the company standardized on a single vendor's platform, more than half of the workforce has access to AI tools, and, by extension, access to agentic AI." The infrastructure is already there.

This built-in availability dramatically reduces barriers to adoption. Employees can begin using agentic features without implementing entirely new systems, which encourages organic experimentation across the organization. As more people experience tangible benefits, adoption accelerates through word of mouth and via demonstration effects — exactly what diffusion theory predicts.

The implication for leaders: Agentic AI's rapid spread isn't an accident. It's happening because the technology is designed to minimize adoption friction. This means that competitive benefits will come not from early access to the technology but from superior organizational design around it.

Organizations and individuals can use agentic AI to differentiate themselves because they can now take advantage of the ways that agentic AI does not fit traditional management frameworks. The same system that offers cost reductions through tool-like automation also enables revenue expansion through workerlike adaptability, innovation acceleration through continuous learning, and organizational learning through human-agent interactions. But accessing these multiple sources of value requires navigating operational tensions that no existing framework addresses.

The path forward requires understanding four operational tensions that expose the inadequacy of traditional management approaches and then redesigning fundamental processes — work design, governance, workforce planning, learning, and investment — to work with rather than against agentic AI's inherent duality. This report offers evidence-based recommendations on how to proceed.

Strategic Tensions When Adopting Agentic AI

The competing pressures executives face aren't abstract theoretical problems. They manifest as specific, irreconcilable conflicts in day-to-day operations. Our research identified four distinct tensions that emerge when organizations try to integrate agentic AI into existing workflows. Each tension represents a fundamental clash between established management principles, forcing leaders to choose to either apply incompatible approaches or develop entirely new, hybrid frameworks. Successfully navigating this challenge requires that leaders manage the following tensions:

1. Scalability versus adaptability. Tools scale predictably; workers adapt dynamically. Agentic AI's ability to do both simultaneously requires new organizational design principles.

2. Experience versus expediency. When is the right time to invest in agentic systems and how should those investments be made? Leaders are faced with balancing long-term capability building with short-term returns.

3. Supervision versus autonomy. How do you supervise something designed to work autonomously? Traditional oversight models assume either full human control or complete automation, not systems that require some human control and differing degrees of automation.

4. Retrofit versus reengineer. When, and by how much, should organizations change processes? That decision requires resources and attention that most change-management frameworks don't address.

The organizations that will succeed are those that recognize agentic AI's dual nature as a feature, not a bug. Strategies that embrace the ambiguity and develop hybrid approaches rather than forcing these systems into existing management categories benefit from both their tool-like scalability and workerlike adaptability.

1. The Flexibility Tension: Scalability Versus Adaptability

Human workers are maximally flexible. They can switch tasks, learn new skills, and adapt to unexpected situations with minimal retraining. Tools are much less flexible. Machinery and infrastructure excel at specific purposes and scale predictably, but they can struggle to adapt to change.

Agentic AI sits in between: more adaptable than tools but (currently) less flexible than workers. How should organizations design processes with intermediate flexibility? Agentic AI's ability to scale predictably and adapt dynamically requires new organizational design principles that don't fit neatly into existing management systems.

In the heart of Goodwill Industries' bustling sorting facility, a river of donations flows ceaselessly with a chaotic mix of vintage T-shirts, designer jeans, and forgotten fabrics. A team of dedicated workers painstakingly sifts through this mountain of textiles, but AI agents are increasingly able to help. This system is not a simple computer vision machine, rigidly programmed to recognize only a few specific items. Instead, it is a flexible AI that constantly learns from and adapts to the ever-changing stream of goods: It can learn to distinguish between a cashmere sweater and a wool blend, identify a rare collectible from a common toy, and even spot subtle signs of wear and tear that might make an item unsuitable for resale. Steve Preston, president and CEO of Goodwill, says, "Our supply chain from beginning to end is very complicated and requires a lot of human intervention, a lot of decision points, and a lot of judgment. So throughout

Our research identified four distinct tensions that emerge when organizations try to integrate agentic AI into existing workflows.

that process, we see a lot of opportunities to incorporate AI in the entire flow of goods, the decision-making process, and making sure that everything we receive finds its best home somewhere in the cycle.”

Goodwill is piloting AI technology that could ultimately be used to sort billions of pounds of donations each year,

ensuring that each item finds its way to the right place, whether it’s a boutique, an online store, or a recycling facility. The AI’s ability to handle complex judgment tasks is encouraging leaders to replace a decades-old, human-centric workflow with new processes. They have begun to design processes for and with agentic AI.




FIGURE 2
AI's Role in the Organization

Our data shows survey respondents believe AI will become more like an assistant, coach, mentor, or colleague rather than a rival or boss.

As the composition of workers shifts, deploying scalable, flexible AI workers creates a strategic tension for process design across the enterprise.

AI systems already act as assistants, colleagues, mentors, coaches, and even as rivals and bosses. Most of our survey respondents anticipate that AI will become more involved in each of these roles within three years. (SEE FIGURE 2, PAGE 6.) As the composition of workers shifts, deploying scalable, flexible AI workers creates a strategic tension for process design across the enterprise. Agentic systems might work best in standardized environments, but over-standardization can eliminate their ability to learn and develop humanlike adaptability — a capability that can help organizations handle edge cases and system failures. Organizations need both the efficiency of standardized processes and the flexibility of humanlike improvisation.

The threat: Organizations that optimize for AI efficiency miss out on AI's humanlike adaptive responses to system failures or unexpected market shifts.

The opportunity: Companies that strike the right balance can achieve both AI-driven efficiency and human-powered adaptability, creating strategic benefits.

2. The Investment Tension: Experience Versus Expediency

Traditional tools require large upfront costs but deliver predictable returns through established depreciation schedules. Human workers are an ongoing variable expense, but their value appreciates with experience and training. Agentic AI defies both models, requiring substantial initial development costs and ongoing variable costs, such as training models on new data. While many technology systems require ongoing maintenance, agentic AI systems simultaneously depreciate through model drift while appreciating through fine-tuning and emergent

capabilities. Should investments in agentic AI be viewed more like investments in tools or in workers, or both?

Directly addressing that question reveals two critical investment tensions that challenge conventional financial planning:

Timing: The moving-target problem. AI systems are evolving quickly, creating uncertainty about when an organization should make significant investments in them. Adopt too early and risk technological obsolescence; wait too long and risk strategic benefits. As Jeff Reihl, executive vice president and technology chairman at LexisNexis Legal & Professional, observes, “This technology is changing so fast, we might have to do a quick catch-up.” Chevron's chief data and analytics officer, Margery Connor, echoes this sentiment, describing the need to remain adaptive to emerging tools and updates: “The fast-paced development of agentic AI requires organizations to be agile while consistently upholding their data and AI governance standards,” she says.

Unlike traditional tools with predictable upgrade cycles, agentic AI requires continuous adaptation and learning. Standard net present value calculations fail when the most valuable applications haven't been conceived yet, and conventional timing models can't account for the speed of technological evolution. Plus, applying conventional replacement schedules risks rapid value decay as systems fall behind the technological curve.

Size: Platforms versus point solutions. Large organizations face a fundamental choice: invest heavily in comprehensive AI platforms or pursue smaller, targeted point solutions. The scale of required investment for each

approach can vary dramatically, making it difficult to accurately gauge upfront and continuing costs.

Platform investments demand substantial upfront commitments with uncertain returns. At Capital One, Prem Natarajan, executive vice president, chief scientist, and head of enterprise AI, describes building “dozens of use cases at scale” from a single, substantial platform investment. The full impact of Capital One’s strategy can be assessed through the lenses of technology exploitation and exploration. Similarly, Walter Sun, senior vice president and global head of AI at SAP, explains that creating “a generative AI hub” allows for full life-cycle management of large language models, whereas building LLMs individually into isolated applications requires the costly integration of legacy systems before any returns materialize. SAP focused on the value to the developer ecosystem and used that to determine what would be a sufficient return from having a platform, Sun notes.

Investing in point solutions offers more predictable costs and measurable returns but risks missing the compound value that emerges when AI capabilities are integrated across business functions. Organizations that choose this path may find themselves managing dozens of disconnected AI tools without the infrastructure to scale or adapt.

Organizations investing solely in one approach, whether human-in-the-loop systems or fully autonomous agents, miss the compound value created by agentic AI’s contributions in multiple decision-making scenarios. Each autonomy level serves different risk tolerances and business contexts, but measuring ROI in isolation obscures the strategic value of a diversified AI portfolio.

Agentic AI’s hybrid features make it difficult to accurately measure investment returns over meaningful time frames or even to assess upfront investment requirements. Companies measuring agentic AI returns through conventional depreciation schedules systematically undervalue the continuous-learning and emergent capabilities these systems generate, failing to account for significant portions of actual value creation.

The threat: Organizations that apply traditional investment frameworks to agentic AI systematically underinvest in continuous learning and adaptation, leading to rapid value decay while the compound returns available from integrated AI ecosystems are missed.

The opportunity: Companies that embrace hybrid investment models and diversified AI portfolios can create compounding returns as their systems learn, adapt, and generate unforeseen capabilities across multiple autonomy levels and business contexts. The window for establishing this strategic advantage is narrowing as competitive pressures intensify and technological evolution accelerates. The question isn’t whether to resolve these tensions but, rather, how quickly organizations adapt their investment frameworks to match the reality of agentic AI.

3. The Control Tension: Supervision Versus Autonomy

Tools are fully owned and controlled, behaving predictably once deployed. Workers must be managed through contracts, incentives, and oversight because humans have autonomy and may pursue divergent goals. Agentic AI requires supervision and management like a worker does

Unlike traditional tools with predictable upgrade cycles, agentic AI requires continuous adaptation and learning.




FIGURE 3
Organizations Anticipate Continued
Increases in AI Autonomy

Survey respondents are two to three times more likely to expect AI systems to work independently from humans and have decision-making authority in three years compared to today.


because its outputs can be unpredictable, even though organizations own it like an autonomous tool. How can organizations design processes to effectively supervise an agent that also works autonomously? (SEE FIGURE 3.)

A process that includes supervision inherently keeps humans involved in decision-making processes. Chevron's Connor notes, "We always have a human in the loop to review and analyze the output so we can determine whether it makes sense or not." Chandra Kapireddy, former head of generative AI, machine learning, and analytics at Truist Bank, says, "If you look at the financial services industry, I don't think there is any use case that is actually

customer-facing, affecting the decisions that we would make without a human in the loop."¹

Despite owning the technology, organizations must treat agentic AI systems with the same oversight typically reserved for human employees.

But at the same time, organizations want the benefits of scale that AI agents can offer while avoiding the bottlenecks that human oversight could create. One leader we spoke with asked, "How do you make sure that you have the right set of controls in place? We're calling it human-in-the-loop, or 'human-out-of-the-loop' for some of our low-risk use cases." The decision to move to human-out-



of-the-loop for lower-risk use cases involves context-dependent assessments of risk. Unlike a simple tool with a fixed function, agentic AI operates within a complex environment where its actions can have unintended consequences. That means organizations must craft policies that set boundaries for decision-making and action to serve as safeguards so that organizational behaviors align with strategic objectives, and negative outcomes are prevented. This is analogous to setting policies and providing oversight for human employees. Sun at SAP describes building a generative AI hub for this purpose: “Through that, we can then connect to different large language models. But we also have the ability, then, to put in the guardrails for all the business applications. We can put in our own analytics. We can put in our own security, our privacy, our compliance, and so on.” The need for these controls highlights the fact that AI is not a fully predictable tool but an agent that managers must guide and constrain.

When a tool malfunctions, it’s a defect. When a coworker makes a mistake, it’s a management and learning opportunity. Managers can apply this coworker mindset to AI. Rebecca Finlay, CEO at the Partnership on AI, argues that companies need to be more transparent about their AI failures: “I would love for us to be in a position where we could share more about things that go wrong, because in that sharing, we learn.” This approach treats AI not as an infallible tool but as an agent whose errors must be understood, managed, and learned from. The learning process parallels how effective teams handle human error.

The threat: Organizations that fail to develop appropriate governance frameworks for agentic AI may face compliance failures, misaligned outputs, or runaway systems that damage business operations.

The opportunity: Companies that master the art of managing artificial colleagues can scale specialized capabilities without the traditional constraints of human hiring, training, and retention.

4. The Scope Tension: Retrofit Versus Reengineer

Agentic AI presents leaders with a critical resource-allocation choice, now complicated by the rapid pace of technological change. The decision is whether to retrofit AI into existing workflows for quick, incremental gains or to reengineer processes for more transformative but slower results.

Retrofitting often requires less time and capital, delivering faster returns with current technology. Reengineering, in contrast, demands a significant commitment of resources over a longer period. This extended timeline creates a new risk: The underlying AI technology may advance significantly before the project is complete, potentially rendering the new process outdated upon launch.

This forces a complex trade-off. Committing to a long-term reengineering project means not only forgoing other, faster optimization projects but also betting on the longevity of today's technology. It may be more strategic to make smaller, incremental improvements now and deliberately skip a technological generation, waiting for a more powerful or stable platform before committing to a complete overhaul. How should leaders decide when to incrementally improve a process versus when to reengineer it entirely around the AI's capabilities?

Goodwill's use of AI for sorting textiles at first appears to be a simple process improvement. However, Preston's comments reveal a much larger ambition. The complexity of the organization's operation, with its billions of unique items and numerous decision points, suggests that AI's potential isn't just in doing a single task faster but in fundamentally changing the entire flow of goods. This presents a dilemma: Should the organization use AI to simply sort better, or to redesign its entire supply chain, from donation to final sale or recycling? Preston says that Goodwill is actively redesigning its operations to better utilize its sorting AI. "The whole process of identifying what the good is, what the best home for that good is, and then how to route it can involve any number of AI applications," he notes. A seemingly incremental application of AI (sorting) opened the door to a complete reengineering of the supply chain following a questioning of the very structure of a decades-old, human-centric workflow.

The threat: Organizations that limit agentic AI to incremental improvements may miss transformative opportunities. Yet a larger redesign can take considerable resources and time.

The opportunity: Companies that question workflows around agentic AI can create new capabilities that differentiate them from competitors.

A Strategic Overhaul of Workflows, Governance, Roles, and Investment

The tensions created by agentic AI's dual nature demand a strategic response that goes beyond incremental adjustment. Unlike technologies that can be managed within traditional functional silos, agentic AI cuts across internal organizational boundaries. For example, it requires IT expertise for technical deployment, HR-like frameworks for performance management, financial models for hybrid investment, legal oversight for autonomous decision-making, and business unit coordination for workflow integration. None of the four fundamental tensions can be resolved by any single function acting alone; each one requires new forms of executive collaboration that transcend the departmental boundaries that have defined organizations' structures since the beginning of industrialization.

The tensions expose not just operational challenges but strategic questions about how organizations create value, structure themselves, govern decision-making, and manage both human and nonhuman contributors over time. Executives are compelled to ask whether agentic AI is intended only to reduce costs or could cover a broader set of values as a source of innovation, differentiation, and organizational learning.

Organizations must treat agentic AI systems with the same oversight typically reserved for human employees.

Their answer will determine how they design processes, structures, governance, and workforce strategies.

Once executives have defined the values they want from agentic AI, whether efficiency, growth, innovation, or some mix, they must address five interlocking implications to move from awareness to action.

1. Redesigning Work: Move Beyond Incrementalism

The rise of agentic AI will reshape not just isolated tasks but entire workflows. For CEOs, the central question is no longer “Where can I automate a step?” but rather “How will process design itself fundamentally change?” Agentic systems don’t simply make existing steps proceed faster; they invite leaders to rethink the design of whole workflows, blending human judgment and machine autonomy in ways that legacy processes were never built to accommodate.

Executives should reflect on the scale of this change. If an organization has 500 processes, how many will be reimaged — 50? 200? 400? Organizations with extensive agentic AI adoption suggest that the answer is far closer to the high end. Our research finds that 66% of organizations with extensive agentic AI adoption expect changes to their

operating model compared with 42% of those that have no plans to adopt. (SEE APPENDIX FIGURE 1, PAGE 19.) This statistic masks a more complex reality: Successful organizations don’t resolve the scope and flexibility tensions. They design workflows that can oscillate between efficiency and adaptability without breaking.

Naomi Lariviere, chief product owner at ADP, offers an example of this approach. Rather than choosing between retrofitting payroll tasks or reengineering the entire process, ADP built its own “agent-building platform” that enables both, simultaneously. Individual payroll tasks can be optimized for efficiency (addressing scalability needs) while the platform architecture allows the rapid deployment of new capabilities as tax rules change across hundreds of locales (addressing adaptability needs). The platform doesn’t eliminate the tension between efficiency and flexibility. It creates infrastructure for managing both demands continuously.

People believe that AI will do twice as much of their jobs in three years. Currently, people feel that AI is doing 23% more of the tasks within their jobs compared with a year ago. However, in three years, people expect that AI will be able to do 46% of their job tasks. Processes will need to change.

People believe that AI will do twice as much of their jobs in three years. Currently, people feel that AI is doing 23% more of the tasks within their jobs compared with a year ago. However, in three years, people expect that AI will be able to do 46% of their job tasks. Processes will need to change.

Similarly, Goodwill's experience reveals how organizations can prepare for scope escalation rather than trying to prevent it. Over time, textile-sorting AI may lead to a major reengineering of the supply chain. Nonetheless, Goodwill is heavily focused on incremental improvements to existing approaches to optimize sorting as the organization looks to a more transformational process with AI.

Almost all organizations adopting agentic AI expect significant process redesign, but organizations further along in agentic AI adoption are more likely (66%) to expect changes in how they are organized and redefine jobs than organizations just beginning adoption (42%). Considering how early all organizations are in agentic AI adoption, this difference is important. Ultimately, agentic AI use is allowing organizations to build new ways to work instead of just adding AI to their old routines.

How to actively manage these competing demands:

- **Build processes with embedded options.** ADP's agent-building platform enables both standardized efficiency and rapid customization across hundreds of locales. Organizations should design workflows that can shift between efficiency and adaptability modes rather than optimizing for one approach.
- **Plan for scope escalation.** Goodwill's textile-sorting AI revealed the need for complete supply chain reengineering. Organizations should establish clear processes for determining when incremental AI improvements should trigger broader redesign discussions, rather than treating scope expansion as project failure.
- **Staff teams with both specialists and orchestrators.** Since our research shows that 43% of extensive adopters anticipate hiring generalists in place of specialists and 45% anticipate a reduction in layers of middle management, organizations should assign teams that include both efficiency optimization specialists and process redesign generalists. (SEE APPENDIX FIGURE 1, PAGES 19 AND 20.)

2. Governance and Decision Rights: Making Decisions and Setting Rules

Agentic AI creates a governance dilemma unlike any previous technology's. Tools are owned and predictable, whereas people are autonomous and must be supervised.

Agentic systems fall somewhere in between: They are owned like assets, but they act in ways that require oversight, akin to employees. The question for managers is no longer "How do we set guardrails for tools?" but "How do we assign decision rights, accountability, and oversight to actors we own but don't fully control?"

Our data shows that 58% of leading agentic AI organizations expect governance structure changes within three years, with expectations that AI systems will have decision-making authority growing 250%. (SEE FIGURE 3, PAGE 9 AND APPENDIX FIGURE 1, PAGE 18.) These organizations aren't solving the supervision-versus-autonomy dilemma. They're creating governance structures that can handle permanent ambiguity about who or what is responsible for making decisions.

The control tension manifests most clearly in how individual organizations simultaneously implement contradictory approaches. Kapireddy describes deploying both human-in-the-loop and human-out-of-the-loop systems at Truist Bank, depending on risk levels, with AI systems managed as both supervised workers (needing human oversight) and autonomous tools (operating independently).² It isn't a strategic choice between approaches; it is an acknowledgment that agentic AI systems require both management philosophies simultaneously.

To provide proper supervision, organizations must recognize that AI agents, like workers, need to be trained and coached. Vibhor Rastogi, head of AI and APAC investments at Citi Ventures, makes this explicit. "As we're looking at startups in the space," he says, "we still feel that these AI agents should be treated like coworkers who need to be trained, coached, supervised. There is some level of supervision, some level of monitoring, some level of oversight that will still have to be provided." Yet, organizations also deploy these systems precisely because they can work autonomously, so a fundamental contradiction results: systems that must be supervised like employees but owned like equipment.

Executives must design governance systems that can apply varying degrees of control depending on the risks, stakes, and context. This is less about finding balance once than it is about building governance capabilities that can continuously adapt.

How to actively manage these competing demands:

- **Clarify decision rights.** Organizations should create governance structures that accommodate both supervision and autonomy. Just as Kapiredy describes using both human-in-the-loop and human-out-of-the-loop for the same AI systems, businesses should establish formal processes where IT, HR, and business units regularly negotiate AI authority boundaries. This approach avoids treating decision rights as static and making permanent control decisions, and acknowledges that autonomy levels will differ by workflows.
- **Build centralized governance infrastructure before deploying autonomous agents.** Since 250% growth is expected in AI decision-making authority, organizations should follow SAP's model of creating governance hubs with enterprisewide guardrails before deploying autonomous systems across business units.

3. Organizational Structure and Strategic Workforce Planning: Redefine Roles, Not Just Skills

Traditional organizational design has been built around humans: spans of control, layers of management, and career paths all focused on human effort. Agentic systems upend this logic. The question for executives is no longer just “What skills do we need?” but “What should the structure of our organization look like when humans and agents are working side by side?”

Our survey data points to the beginnings of structural change. As noted, among organizations with extensive agentic AI adoption, 45% expect reductions in middle management

layers. This isn't simply a workforce adjustment; it implies broader organizational redesign. If agents coordinate workflows, traditional managerial spans of control increase and the number of hierarchical layers decreases. The net effect is flatter organizations where fewer people manage more workers, with human managers increasingly responsible for orchestrating hybrid teams of humans and agents.

Agentic adoption also reshapes workforce composition. When agents learn and adapt across functions, demand rises for generalists who can span domains and work across systems. When agents take on routine tasks, organizations need fewer entry-level employees. This isn't traditional workforce optimization. It is a recognition that organizations are now managing both people and nonhuman contributors, requiring a recalibration of hiring, roles, and career paths.

The data on AI roles reflects this complex scenario. Survey respondents expect AI to become more like assistants or colleagues and mentors or coaches over the next three years. (SEE FIGURE 2, PAGE 6.) This role evolution doesn't resolve the question of whether AI is a worker or tool. It multiplies the identity categories organizations must manage simultaneously. These developments will challenge existing notions of career progression. Promotion and compensation systems built around routine analysis or coordination may become obsolete, forcing organizations to reward new capabilities: orchestrating hybrid teams, managing exceptions, and providing ethical oversight, for example. In this future, “generalist” does not mean junior or inexperienced; it describes leaders who can span domains, manage ambiguity, and supervise human-AI collaboration at scale.

Traditional organizational design has been built around humans. Agentic systems upend this logic.

At organizations with extensive agentic AI adoption, 95% of respondents indicate AI has positively impacted their job satisfaction, suggesting that embracing rather than resolving this identity confusion creates better outcomes than trying to narrowly categorize AI.

How to actively manage these competing demands:

- **Reshape organizational structures.** The organizational shifts expected point to a need to plan for flatter hierarchies and wider spans of control and redefine managerial roles to focus less on information brokering and more on orchestrating hybrid human-AI teams.
- **Create dual career paths in workforce pipelines.** Of the organizations with extensive AI adoption, 43% plan to hire more generalists and 29% expect fewer entry-level roles. (SEE APPENDIX FIGURE 1, PAGE 20.) This suggests that organizations should establish distinct advancement paths for both AI-augmented specialists and AI orchestrators, acknowledging that both roles are essential.
- **Plan for AI's evolving role in teams.** Organizations expect AI to increasingly act as a coach or mentor. To prepare, they can build workforce models and leadership development programs that anticipate these role shifts to ensure that humans are ready to collaborate with agents in more guidance- and judgment-oriented roles.

4. Upskilling, Learning Loops, and Life-Cycle Management: Building Human and Agent Capacity

Agentic AI use reshapes how organizations must think about learning and development on two fronts. First, employees require upskilling, not only in the technical use of AI systems but in the supervision, critique, and orchestration of them. Second, the agents themselves need life-cycle management (onboarding, training, retraining, and eventual retirement) if they are to remain reliable and valuable. Executives should continually ask themselves, “How are we ensuring that our people and agents are continuously learning?”

On the human side, traditional reskilling has focused on teaching people to operate tools. But working alongside agents demands broader capabilities: Employees must know when to trust outputs, when to challenge them, and how to integrate them into team workflows.

On the agentic side, these systems can become more valuable with experience, creating a compounding effect if properly managed. Organizations need to treat agents not as static tools but as dynamic participants with life cycles of their own that require structured support similar to that afforded to human employees.

BLENDING TOOL AND WORKER IN A SINGLE SYSTEM

Sun at SAP sees agentic AI developing an expertise that affords a human-worker level of expertise to handle complex tasks. “The way I see it is that if you had a choice of accessing a very large encyclopedia or the alternative of a person who’s an expert in a particular task, which one would you pick? We feel that for a lot of our business tasks, we need that expert,” he explains. “And that’s why we feel that by being able to fine-tune and create these smaller language models, we get that expertise built into an AI agent.” The initial model (the encyclopedia) is a depreciating fixed asset, but through the ongoing process of fine-tuning, it becomes a more valuable, flexible “expert.”

To augment entrepreneurs’ expertise with expertise in areas outside their core strengths, Alibaba.com developed a four-agent system (product listing, communication, marketing, and risk management) for small businesses. But the agents aren’t tied to a fixed way of working. Instead, company president Kuo Zhang notes, “each of the agents can beat 90% of the real humans doing this job. The more you work with the agent, the more it understands your business.” The development initially required tool-like design but retained benefits normally associated with human workers. With increased experience, the agents are learning more and more about specific situations. As a result, small entrepreneurs can increasingly use the agents to handle problems they face in their own contexts.

Connor notes that at Chevron, model updates arrive “every couple of months,” underscoring the need for ongoing vigilance to keep systems current. Natarajan points to continuous reinvestment in fine-tuning at Capital One, to ensure that agents don’t drift off course. If organizations fail to actively care for their agents, the systems will quickly become outdated, inaccurate, or even risky.

Just as HR teams recruit, onboard, train, evaluate, and eventually retire employees, organizations adopting agentic AI need parallel support functions for their nonhuman workforce. These functions include onboarding processes (testing and validating new agents), performance reviews (tracking agents’ accuracy, adaptability, and bias), and retraining (fine-tuning agents as new data becomes available). Some companies are experimenting with assigning explicit accountability for this. Moderna, for instance, merged its tech and HR departments, thus making it clear that agents must be managed as part of the workforce, not just IT infrastructure.³

How to actively manage these competing demands:

- **Upskill humans to orchestrate, not just operate.** Organizations must train employees to supervise, redirect, and critique agent outputs, not just consume them.
- **Redesign onboarding.** Agents should be embedded into training programs so new hires learn to work with them from day 1.
- **Create an “HR for agents” function.** Organizations must assign responsibility for recruiting, onboarding, evaluating, retraining, and retiring AI agents, just as HR provides those functions for humans.

5. Investment Strategy: Budgeting for Permanent Uncertainty

Traditional investment frameworks assume that organizations can choose between competing approaches: capital versus operational expenses, short-term versus long-term returns, centralized versus distributed spending. But because agentic investments cut across these categories, organizations need financial architectures that can accommodate multiple investment approaches.

Agentic AI’s swift evolution requires adaptability. With the rapid emergence of new models and features, investment decisions must anticipate switching costs and plan

for continuous reinvestment in systems. Financial models should be flexible enough to account for rapid upgrades, retraining, and replacement, without locking organizations into static depreciation schedules.

Unlike most capital investments, agentic AI can appreciate with use. As managers fine-tune agents, expose them to more data, and embed them in workflows, agentic systems can become more effective. While most technologies plateau or decline in usefulness over time, agentic systems can gain value the more they are used, in spite of possible model drift or technology changes. Organizations can use these systems to become more proactive, spot market signals earlier, and compete on differentiation rather than just efficiency. This appreciation requires a strategic departure from traditional investment logic.

Organizations should anchor investments in customer and value perspectives. The question is not simply how to fund AI but what value it is meant to deliver — whether that is cost efficiency, faster innovation, or deeper customer engagement. Our research shows that organizations further along in adoption are much more likely to expect fundamental changes in how they differentiate competitively. This demands that their investment choices be directly tied to strategic goals rather than dispersed across a patchwork of disconnected projects.

How to actively manage these competing demands:

- **Build for adaptability.** Organizations must assume that there will be rapid technology shifts and therefore should design budgets and systems that can incorporate switching costs and support continuous retraining.
- **Create investment review processes that accommodate functional disagreement/joint alignment across businesses.** Organizations should establish monthly review cycles where IT, finance, HR, and business units can advocate for contradictory investment approaches without requiring consensus.
- **Develop metrics that track competing sources of value creation.** Organizations should measure AI systems as appreciating assets (learning capabilities, expanded applications) and depreciating tools (model drift, technological obsolescence) simultaneously and accept that such metrics will sometimes point in opposite directions.

Conclusion

With agentic AI, leaders are managing a new entity without historical precedent. It is a tool that learns, a worker that is owned, and an investment that behaves like both tool and worker. Agentic AI does not fit neatly within the substitute-or-complement framework because it demands to be managed as both a worker and a tool at once. The substitute-complement framework asks, “Is AI replacing coworkers or augmenting them?” But agentic AI forces a more unsettling question: “How do we manage artificial colleagues that we own like equipment but must supervise like people, and that depreciate like machinery but learn like humans?” This new technology creates critical design challenges.

Companies beginning to use agentic AI are already grappling with that question. Some are treating it primarily as a way to reduce costs whereas others are using it to expand capacity, speed up experimentation, or elevate the baseline performance of their people. Leaders need to be explicit about what they are optimizing for and design their processes, structures, and governance accordingly.

Treating an agentic AI deployment like a traditional tool misses its flexibility advantages, but treating it like a newly hired staff member underestimates the infrastructure requirements. This dual risk creates unavoidable tensions between flexibility and efficiency, control and autonomy, and upfront investment and continuous learning. Success will be defined not by eliminating these conflicts but by mastering them.

The challenge of agentic AI is organizational, not technological. Many organizations are adopting this technology at a breakneck pace, often before they have a coherent strategy in place. But the speed of adoption is not a measure of progress. Advantage will not come from having the best AI injected into the organization as quickly as possible but from having the best answer to the question, “How do we reorganize our organization around it?” The “adopt now, strategize later” approach is a high-stakes gamble that avoids the hard work of identifying a purpose for agentic AI that fits well with the organization’s strategy.

Ultimately, this challenge is deeply human and requires that businesses break down organizational silos. The era of managing technology solely within the IT department is over; governance is now a mandatory, cross-functional effort where IT, HR, finance, and operations must collaborate on

a unified framework. Agentic AI elevates human judgment rather than eliminating it. Strategic oversight, ethical governance, and the ability to orchestrate human-AI teams become the most critical human skills as AI agents handle tasks previously performed by human workers. The organizations that thrive will be those that focus less on the technology itself and more on the human systems that surround it. This forces a deeply unsettling question for today’s leaders: “Are we simply adding a new tool to our business, or are we introducing a new, nonhuman actor into our organization?” How we respond will define the next era of management.


APPENDIX

1. Use of Agentic AI and Organizational Change

As organizations increase their adoption of agentic AI, they expect significant and widespread changes across the business. Our data reveals a consistent, direct relationship: The greater the level of agentic AI adoption, the higher the percentage of organizations that anticipate fundamental shifts in decision-making, structure, talent strategy, and competitive differentiation.


Decision-making: How does the use of agentic AI relate to the processes, accountability, and rules that guide how an organization makes choices? Compared with organizations that have no plans to adopt agentic AI, the companies that are already making extensive use of the technology foresee bigger changes in how they will make decisions. For example, the expectation of a change in governance structure in three years is 37% among those with no plans but 58% among those already using agentic AI extensively. Similarly, the belief that the time horizon for decisions is changing is 42% versus 46%, respectively.

Our organization's use of AI leads to ...




APPENDIX FIGURE 1
How AI Will Reshape Organizations

Respondents at different stages of agentic AI adoption were asked whether they believe the technology will shape decision-making, organizational structure, talent strategy, and competitive differentiation at their company.



Organizational structure: How does the use of agentic AI relate to an organization's hierarchy, roles, and overall operating model? The data links higher agentic AI adoption to a greater anticipation of structural reorganization. The expectation of a change in the operating model and a redefinition of roles is 42% in organizations with no plans to adopt agentic AI compared with 66% among those using it extensively. Likewise, the percentage of organizations that expect to see a reduction in the layers of middle management is 30% versus 45%, respectively.

In 3 years, our organization's use of AI will lead to ...




Talent strategy: With increasing use of agentic AI, how will organizations have to change how they hire, develop, and manage talent? Increased AI adoption positively correlates with the expectation of shifts in how companies manage their workforces. In organizations with increased agentic AI adoption, people are more likely to anticipate a move toward hiring generalists over specialists (28% among those with no plans and 43% among those with extensive plans).

In 3 years, our organization's use of AI will lead to ...



Competitive differentiation: If everyone uses agentic AI, how will organizations and individuals stand out? This category shows the most pronounced differences in expectations. Organizations overwhelmingly believe that as AI adoption becomes extensive, it will alter how both companies and individuals stand out. The expectation that AI will change an organization's ability to differentiate itself is 53% for those with no agentic AI adoption plans compared with 73% among those using it extensively. Most significantly, the belief that AI will affect an individual's ability to differentiate themselves from coworkers shows the largest gap, at 49% versus 76%.

Our organization's use of AI leads to ...



2. Working With AI Versus Working With Humans


Our research compares how working with humans, generative AI, and traditional AI affects five performance attributes. Both AI types are superior to human coworkers for enhancing performance metrics. For instance, 90% of respondents reported an increase in speed when working with generative AI, and 89% saw a boost in efficiency, compared with only 48% and 58%, respectively, from collaborating with humans. More notably, respondents believe that generative AI also outperforms humans in stimulating certain cognitive traits, with 79% believing that GenAI boosts people's curiosity (versus 71% feeling that way about other humans) and 68% believing that GenAI increases people's agency (versus 57% for humans). While humans maintain an edge in fostering creativity (77%), generative AI is a close second at 76%, significantly higher than traditional AI's 58%. Overall, the data suggests that while traditional AI is valued primarily for speed and efficiency, respondents view generative AI as a partner that enhances both performance and human-centric qualities like curiosity and agency.

Conversely, the data on negative effects of each collaboration type for each attribute reveals the potential downsides. Working with humans is viewed as most likely to decrease speed (by 22% of respondents) and efficiency

(by 16%), reinforcing the idea that AI collaboration is well suited to performance-based tasks. In contrast, both generative AI and traditional AI are seen as highly unlikely to reduce speed or efficiency, with expectations of decreases at only 2%-3%. However, traditional AI is seen as the most likely collaborator to stifle creativity (by 13% of respondents) and curiosity (by 10%). This suggests a trade-off: While human collaboration may be slower, it's least likely to dampen creativity, whereas traditional AI boosts performance but at a greater risk to creative and curious exploration.

Overall, our findings also indicate a strong positive correlation between the adoption of agentic AI and employee job satisfaction. Organizations with the most extensive AI adoption report the highest levels of satisfaction, with 95% of individuals at organizations with the most extensive agentic AI adoption reporting AI positively impacting their job satisfaction. This correlation suggests that as companies more deeply integrate sophisticated, autonomous AI, it can lead to more fulfilling work and higher job satisfaction, possibly by automating tedious tasks and allowing employees to focus on more strategic activities.

For people in my functional area, how does working with AI compare to working with humans?



APPENDIX FIGURE 2
How Working With AI Affects
Various Human Competencies


Working with AI primarily increases speed and efficiency.






3. Agentic AI Adoption and Decision-Making

Organizations with the highest level of agentic AI adoption are far more likely to see increased applications of the technology across various modes of decision-making. But the largest impact of agentic adoption is augmenting human judgment: Seventy-nine percent of extensive agentic AI adopters say they are investing in using AI to generate insights for a human decision maker. In contrast, fully autonomous scenarios where “AI decides and implements” are significantly less common, peaking at 54% in the same group. This indicates that experienced organizations are currently primarily investing in using AI as an analytical partner to augment human oversight rather than as an independent decision maker.



APPENDIX FIGURE 3 Investment in Five Human-AI Interaction Modes


The largest investment in agentic AI is to augment human judgment.



4. Hopes and Fears About Task Replacement

Do people hope or fear that AI will take over some of their job tasks within five years? The data shows that as organizations adopt more capable AI, people's feelings of hope and fear do not significantly change. Across all stages of agentic AI adoption, from companies with no plans to those with extensive implementations, the hope that AI will handle certain tasks remains high, fluctuating between

78% and 85%. At the same time, fear about AI performing these tasks stays relatively low, ranging from 21% to 32%. This pattern is consistent with a longer-term trend: Hope increased from 70% in 2017 to 84% in 2024, while fear decreased from 31% in 2017 to 20% in 2024. Ultimately, greater exposure to AI in the workplace does not appear to substantially alter these sentiments.



APPENDIX FIGURE 4 AI Hope Continues to Outpace Fear


Organizations at all levels of agentic AI adoption are more hopeful than fearful that AI will automate elements of their work.

5. Authenticity

When organizations replace human tasks and interactions with automated computer systems, they risk appearing more robotic and less authentic. Human communication is filled with nuance, empathy, and spontaneity — qualities that automated systems struggle to replicate. A chatbot or an AI-generated email, for example, can feel impersonal and scripted, lacking the genuine connection of a conversation with a real person. As a result, agentic AI adoption could make an organization seem distant and cold, particularly when handling sensitive customer or employee issues, where empathy is crucial. Overreliance on automation can erode the human touch — the qualities that build trust and make an organization seem authentic and relatable.

Our findings indicate some support for this concern. There is a tension between using AI for efficiency and maintaining a sense of genuine human ownership and decision-making. There's also a notable disconnect between how people present their AI-assisted work and their willingness to

disclose its use. Fifty percent of respondents believe that their AI-assisted performance is viewed as entirely their own and 51% report that they let others know when they've used AI. This suggests that many people are comfortable taking full credit for AI's contributions, which raises questions about workplace transparency and what it means to be authentic. Authenticity also varies with power dynamics. Perceptions of inauthenticity change depending on who is using the AI. Only 35% of respondents feel that their leaders become less authentic when using AI to communicate. Meanwhile, 23% of respondents consider their own direct reports to be less authentic when they use AI to make decisions. These figures indicate that employees are less skeptical of AI's role in critical decision-making, especially by subordinates, than they are of its use in top-down communication. It highlights that the type of task being automated heavily influences whether the outcome is perceived as authentic.



APPENDIX FIGURE 5

How Using AI Affects Authenticity

Respondents are generally not concerned about a lack of authenticity when using AI.

6. The Emerging Agent-to-Agent Ecosystem

Originally, managers designed companies around human interaction. Increasingly, they are preparing for autonomous AI agents to interact directly with one another, a trend that correlates strongly with agentic AI adoption. Our research shows that as organizations move beyond initial trials, they prioritize enabling this capability. For example, while we find that only 30% of companies in the agentic AI pilot stage enable internal agent-to-agent interaction, that figure is 52% for organizations with extensive agentic

AI adoption. This pattern also holds for external processes. The readiness to allow AI agents to interact with outside partners grows significantly with adoption levels, at just 17% in the pilot phase but 27% among extensive adopters. This difference underscores that as companies integrate AI more deeply into their operations, they see a greater need for agents to autonomously manage tasks like negotiating with suppliers or coordinating logistics.




APPENDIX FIGURE 6 Agentic AI Adopters Are More Ready for Agent-to-Agent Interaction

While organizational readiness for fully autonomous AI agents is still nascent, a larger percentage of organizations with extensive agentic AI adoption are ready for their agents to interact externally.

7. AI Will Do More Current Tasks Than New Tasks

An organization's adoption of agentic AI directly correlates with what employees expect AI to do in the next three years. As companies advance from having no plans for agentic AI to extensive implementation, employees increasingly anticipate that AI will perform both their current tasks and new tasks. For example, the expectation that AI will do current tasks is 15% in organizations with

no plans versus 49% in those with extensive adoption. There is a similar difference for new tasks, at 11% and 39%, respectively. At the highest level of adoption, a notable gap emerges, as significantly more people believe that AI will take over their current work (49%) versus performing new kinds of work (39%). People may find it easier to envision AI performing familiar tasks rather than unfamiliar ones.



APPENDIX FIGURE 7 Agentic AI Adopters Believe AI Will Do More Current Tasks Than New Tasks

An organization's adoption of agentic AI directly correlates with how much of both current and new tasks employees expect AI to do three years from now.

ACKNOWLEDGMENTS

We thank each of the following individuals, who were interviewed for this report:

[Aparna Chennapragada](#)

chief product officer, Microsoft Experiences and Devices, Microsoft

[Margery Connor](#)

chief data and analytics officer, Chevron

[Rebecca Finlay](#)

CEO, Partnership on AI

[Naomi Lariviere](#)

chief product owner, ADP

[Prem Natarajan](#)

executive vice president, chief scientist, and head of enterprise AI, Capital One

[Steve Preston](#)

CEO, Goodwill Industries International

[Huiming Qu](#)

vice president, data science, The Home Depot

[Vibhor Rastogi](#)

head of AI and APAC investments, Citi Ventures

[Jeff Reihl](#)

executive vice president and technology chairman, LexisNexis Legal & Professional

[Walter Sun](#)

senior vice president and global head of AI, SAP

[Kuo Zhang](#)

president, Alibaba.com

REFERENCES

- 1 S. Ransbotham and S. Khodabandeh, hosts, “Overcoming AI Hallucinations: Truist’s Chandra Kapireddy,” “Me, Myself, and AI,” podcast, season 11, episode 3 (April 15, 2025), *MIT Sloan Management Review*, <https://sloanreview.mit.edu>.
- 2 Ibid.
- 3 I. Bousquette, “Why Moderna Merged Its Tech and HR Departments,” *The Wall Street Journal*, May 12, 2005, www.wsj.com.
- i C. Halton, “Diffusion of Innovations Theory: Definition and Examples,” Investopedia, May 27, 2025, www.investopedia.com.

MIT SLOAN MANAGEMENT REVIEW

At *MIT Sloan Management Review*, we explore how leadership and management are transforming in a disruptive world. We help thoughtful leaders capture the exciting opportunities — and face down the challenges — created as technological, societal, and environmental forces reshape how organizations operate, compete, and create value.

MIT SLOAN MANAGEMENT REVIEW BIG IDEAS

MIT Sloan Management Review's Big Ideas Initiatives develop innovative, original research on the issues transforming our fast-changing business environment. We conduct global surveys and in-depth interviews with front-line leaders working at a range of companies, from Silicon Valley startups to multinational organizations, to deepen our understanding of changing paradigms and their influence on how people work and lead.

BOSTON CONSULTING GROUP

Boston Consulting Group (BCG) partners with leaders in business and society to tackle their most important challenges and capture their greatest opportunities. BCG was the pioneer in business strategy when it was founded in 1963. Today, we work closely with clients to embrace a transformational approach aimed at benefiting all stakeholders — empowering organizations to grow, build sustainable competitive advantage, and drive positive societal impact. Our diverse, global teams bring deep industry and functional expertise and a range of perspectives that question the status quo and spark change. BCG delivers solutions through leading-edge management consulting, technology and design, and corporate and digital ventures. We work in a uniquely collaborative model across the firm and throughout all levels of the client organization, fueled by the goal of helping our clients thrive and enabling them to make the world a better place.

BCG HENDERSON INSTITUTE

The BCG Henderson Institute is Boston Consulting Group's strategy think tank, dedicated to exploring and developing valuable new insights from business, technology, science, and economics by embracing the powerful technology of ideas. The Institute engages leaders in provocative discussion and experimentation to expand the boundaries of business theory and practice and to translate innovative ideas from within and beyond business. For more ideas and inspiration from the Institute, please visit bcghendersoninstitute.com.

BCG X

BCG X is the tech build and design unit of BCG.

Turbocharging BCG's deep industry and functional expertise, BCG X brings together advanced tech knowledge and ambitious entrepreneurship to help organizations enable innovation at scale.

With nearly 3,000 technologists, scientists, programmers, engineers, and human-centered designers located across 80-plus cities, BCG X builds and designs platforms and software to address the world's most important challenges and opportunities.

Teaming across our practices, and in close collaboration with our clients, our end-to-end global team unlocks new possibilities. Together we're creating the bold and disruptive products, services, and businesses of tomorrow.

Get more on artificial intelligence from *MIT Sloan Management Review*:

Read the report online at <https://sloanreview.mit.edu/ai2025>

Visit our site at <https://sloanreview.mit.edu/big-ideas/artificial-intelligence-business-strategy>


Get free newsletter updates at <https://sloanreview.mit.edu/offers-enews>

Reproducing or distributing this report requires written permission. To request permission, visit our website at shop.sloanreview.mit.edu or email smr-help@mit.edu





sloanreview.mit.edu/big-ideas



MIT Sloan
Management Review

[SLOANREVIEW.MIT.EDU/BIG-IDEAS](https://sloanreview.mit.edu/big-ideas)